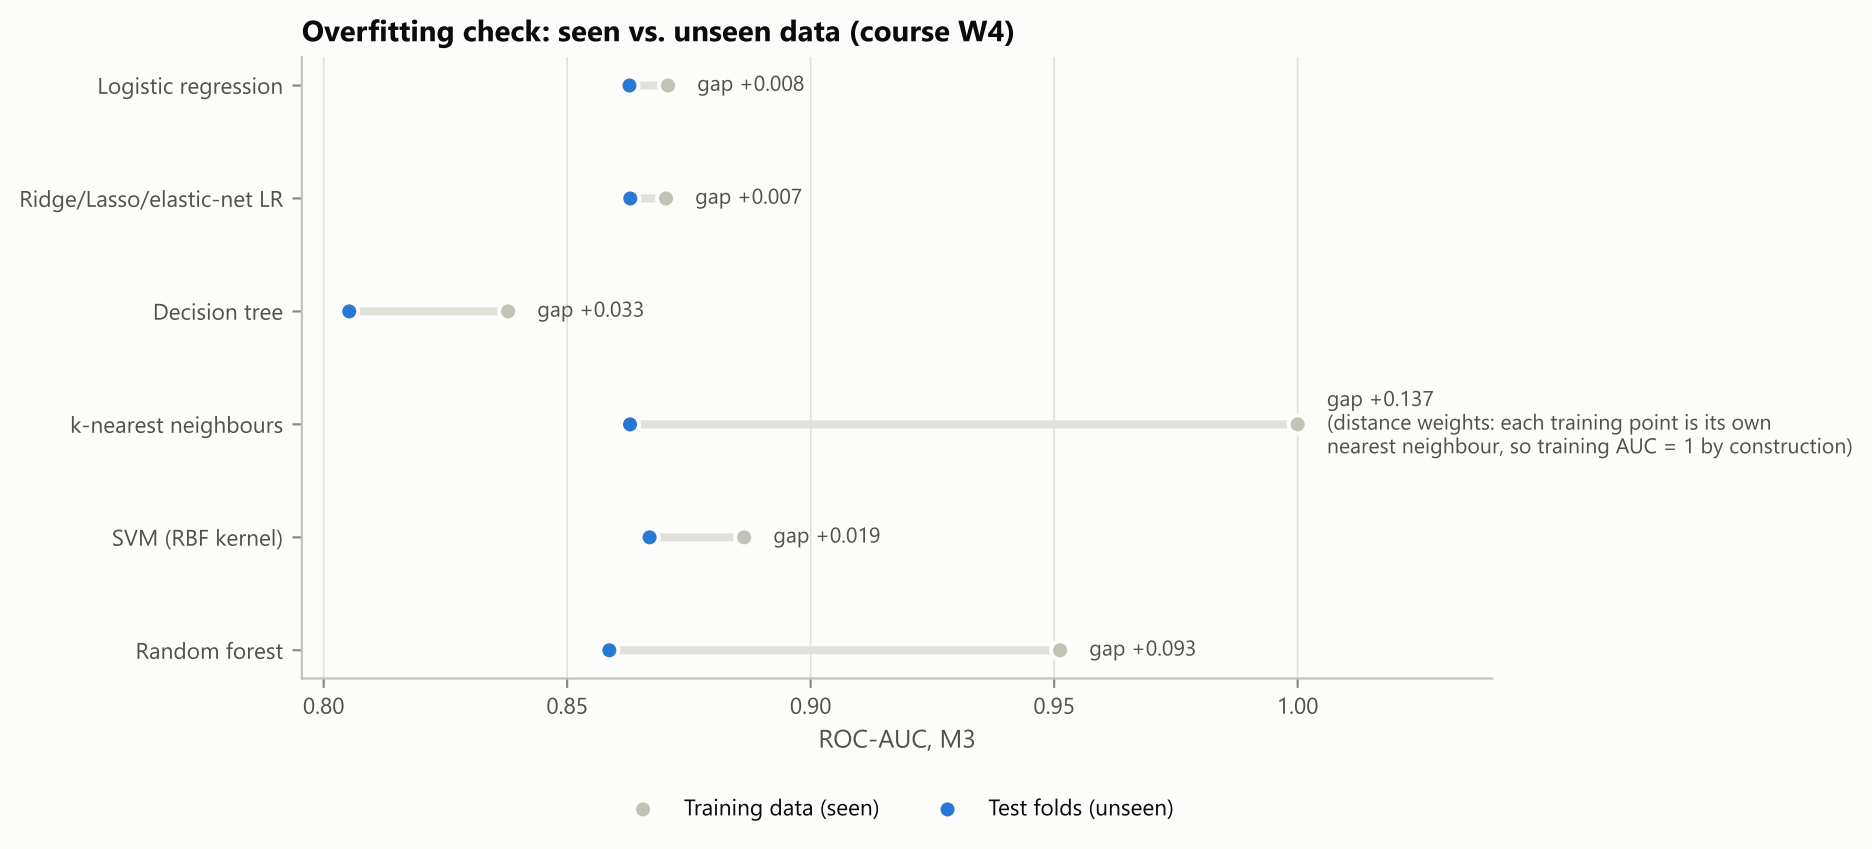
Overfitting check: seen vs. unseen data (course W4)
Logistic regression
Ridge/Lasso/elastic-net LR
Decision tree
k-nearest neighbours
SVM (RBF kernel)
Random forest
gap +0.008
gap +0.007
gap +0.033
gap +0.137
(distance weights: each training point is its own
nearest neighbour, so training AUC = 1 by construction)
gap +0.019
gap +0.093
0.80
0.85
0.90
0.95
1.00
ROC-AUC, M3
Training data (seen)
Test folds (unseen)